

Demonstrate-the-use-ridge-regression-to-predict-the-mileage-of-the-car-using-mtcars-dataset.R

Harini G

```
#Name: Harini G
```

```
#1.Load mtcars dataset
```

```
data(mtcars)
```

```
#?mtcars
```

```
#2.install ridge and glmnet packages
```

```
#install.packages("glmnet")
```

```
#install.packages("ridge")
```

```
#3.Perform the exploratory data analysis
```

```
nrow(mtcars)
```

```
## [1] 32
```

```
ncol(mtcars)
```

```
## [1] 11
```

```
#inference:the dataset contains 32 rows and 11 coulmns
```

```
head(mtcars)
```

```
##           mpg cyl  disp  hp  drat    wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160  110  3.90  2.620  16.46  0  1    4    4
## Mazda RX4 Wag  21.0   6  160  110  3.90  2.875  17.02  0  1    4    4
## Datsun 710      22.8   4  108   93  3.85  2.320  18.61  1  1    4    1
## Hornet 4 Drive  21.4   6  258  110  3.08  3.215  19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360  175  3.15  3.440  17.02  0  0    3    2
## Valiant         18.1   6  225  105  2.76  3.460  20.22  1  0    3    1
```


```
tail(mtcars)
```

```
##           mpg cyl  disp  hp  drat    wt  qsec vs am gear carb
## Porsche 914-2  26.0   4 120.3   91  4.43  2.140  16.7   0  1    5    2
## Lotus Europa   30.4   4  95.1  113  3.77  1.513  16.9   1  1    5    2
## Ford Pantera L  15.8   8 351.0  264  4.22  3.170  14.5   0  1    5    4
## Ferrari Dino    19.7   6 145.0  175  3.62  2.770  15.5   0  1    5    6
## Maserati Bora   15.0   8 301.0  335  3.54  3.570  14.6   0  1    5    8
## Volvo 142E      21.4   4 121.0  109  4.11  2.780  18.6   1  1    4    2
```


```
summary(mtcars)
```

```
##      mpg      cyl      disp      hp
## Min.   :10.40   Min.   :4.000   Min.    : 71.1   Min.    : 52.0
## 1st Qu.:15.43   1st Qu.:4.000   1st Qu.:120.8   1st Qu.: 96.5
## Median :19.20   Median :6.000   Median :196.3   Median :123.0
## Mean   :20.09   Mean   :6.188   Mean    :230.7   Mean    :146.7
## 3rd Qu.:22.80   3rd Qu.:8.000   3rd Qu.:326.0   3rd Qu.:180.0
## Max.   :33.90   Max.   :8.000   Max.    :472.0   Max.    :335.0
##      drat      wt      qsec      vs
## Min.   :2.760   Min.   :1.513   Min.    :14.50   Min.    :0.0000
## 1st Qu.:3.080   1st Qu.:2.581   1st Qu.:16.89   1st Qu.:0.0000
## Median :3.695   Median :3.325   Median :17.71   Median :0.0000
## Mean   :3.597   Mean   :3.217   Mean    :17.85   Mean    :0.4375
## 3rd Qu.:3.920   3rd Qu.:3.610   3rd Qu.:18.90   3rd Qu.:1.0000
## Max.   :4.930   Max.   :5.424   Max.    :22.90   Max.    :1.0000
##      am      gear      carb
## Min.   :0.0000   Min.   :3.000   Min.    :1.000
## 1st Qu.:0.0000   1st Qu.:3.000   1st Qu.:2.000
## Median :0.0000   Median :4.000   Median :2.000
## Mean   :0.4062   Mean   :3.688   Mean    :2.812
## 3rd Qu.:1.0000   3rd Qu.:4.000   3rd Qu.:4.000
## Max.   :1.0000   Max.   :5.000   Max.    :8.000
```


```
#Histogram
hist(mtcars$mpg)
```




```
#inference: miles/ gallon is left skewed
hist(mtcars$disp)
```




#inference: the displacement is having Multi-Modal Distribution
`hist(mtcars$hp)`




```
#inference: gross horsepower is left skewed  
hist(mtcars$drat)
```




```
#inference: Rear axle ratio is following normal distribution  
hist(mtcars$wt)
```



#inference: weight is having The Bi-Modal Distribution
`hist(mtcars$qsec)`




```
#inference: mile time is following normal distribution  
hist(mtcars$carb)
```




```
#inference: carburetors is having The Bi-Modal Distribution
```


```
#boxplot  
boxplot(mtcars$mpg, main="mpg")
```




```
boxplot(mtcars$disp, main="disp")
```




```
boxplot(mtcars$hp, main="hp")
```



```
boxplot(mtcars$drat, main="drat")
```



```
boxplot(mtcars$wt, main="wt")
```

```
boxplot(mtcars$qsec, main="qsec")
```

#4. Choose optimum lambda value

```
x_var=model.matrix(mpg~.,mtcars)[,-1]
```

```
y_var=mtcars$mpg
```

```
lambda_seq =10^seq(2, -2, by = -.1)
```

```
set.seed(86)
```

```
train = sort(sample(1:nrow(x_var),0.8*nrow(x_var)))
```

```
x_test = (-train)
```


```
y_test = y_var[x_test]
```

#5. Extract the model using k-cross validation

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1
```



```
cv_output = cv.glmnet(x_var[train,], y_var[train], alpha = 1, lambda =
lambda_seq)
```

```
## Warning: Option grouped=FALSE enforced in cv.glmnet, since < 3
observations per
## fold
```

```
best_lambda= cv_output$lambda.min
best_lambda
```

```
## [1] 0.3981072
```

```
best_fit=cv_output$glmnet.fit
summary(best_fit)
```

```
##           Length Class      Mode
## a0          41    -none-  numeric
## beta        410   dgMatrix S4
## df           41    -none-  numeric
## dim           2    -none-  numeric
## lambda       41    -none-  numeric
## dev.ratio    41    -none-  numeric
## nulldev       1    -none-  numeric
## npasses       1    -none-  numeric
## jerr          1    -none-  numeric
## offset        1    -none-  logical
## call          5    -none-  call
## nobs          1    -none-  numeric
```

#6. Build the final model and interpret

```
library(ridge)
inputData = data.frame (mtcars)
trainingData=inputData[train, ]
testData = inputData[-train, ]
linRidgeMod =linearRidge(y_var ~ x_var, data = trainingData)
linRidgeMod

##
## Call:
## linearRidge(formula = y_var ~ x_var, data = trainingData)
##
## (Intercept)      x_varcyl      x_vardisp      x_varhp      x_vardrat
## 19.228871873 -0.244604033 -0.001631983 -0.013126081  0.971434008 -
## 1.935291447
##      x_varqsec      x_varvs      x_varam      x_vargear      x_varcarb
##  0.327054947  0.469354207  2.131844699  0.639243640 -0.660066576

predicted=predict(linRidgeMod, testData) # predict on test data

## Warning: 'newdata' had 7 rows but variables found have 32 rows

compare = cbind (actual=testData$response, predicted) # combine
compare

##      predicted
## 1      22.20627
## 2      21.89592
## 3      26.68822
## 4      20.75135
## 5      16.95415
## 6      20.34093
## 7      14.13596
## 8      22.85553
## 9      23.68130
## 10     19.39501
## 11     19.59124
## 12     15.19320
## 13     15.91661
## 14     15.95067
## 15     11.65483
## 16     11.22209
## 17     11.30333
## 18     27.82737
## 19     29.00313
## 20     28.83592
## 21     23.87643
## 22     16.76810
## 23     17.47494
## 24     13.99427
```

```
## 25 16.04689
## 26 28.15331
## 27 26.49133
## 28 27.35073
## 29 18.62860
## 30 19.82039
## 31 13.73599
## 32 25.15599
```

```
mean (apply(compare, 1, min)/apply(compare, 1, max))
```

```
## [1] 1
```

```
#accuracy is 1
```